

Solving Maximum Weight Independent Set in Practice Using Tree Decompositions

BJARN DE JONG

We present a C++ implementation of an exact solver for the MAXIMUM WEIGHT INDEPENDENT SET problem using dynamic programming over tree decompositions. Given a graph G and a tree decomposition of width $k - 1$, the algorithm runs in time $O(n \cdot k \cdot 2k \cdot 2^k)$, where n is the number of vertices in the graph. The dominant cost arises from enumerating independent sets within each bag, which in the worst case can be as many as 2^k per bag. However, in practice, the number of independent sets encountered is typically much smaller. Our implementation reliably solves instances of graphs with tree decompositions for which the maximum bag size k is less than 32, where memory usage remains practical. Although our encodings support tree decompositions with maximum bag size k less than 128 as a data type constraint, memory demands generally limit effective computation to lower widths. Experimental evaluation confirms the solver's effectiveness on graphs with tree decompositions of low width.

1 Introduction

In this work, we give an approach for exactly solving the MAXIMUM WEIGHT INDEPENDENT SET problem. Given an undirected graph $G = (V, E)$, where V is its set of vertices and E is its set of edges, a subset of vertices $S \subseteq V$ is an *independent set* if the elements of S are pairwise nonadjacent in G . If $w(v)$ denotes the weight of vertex $v \in V$, the weight of an independent set S is $w(S) = \sum_{v \in S} w(v)$. In the MAXIMUM WEIGHT INDEPENDENT SET problem (MWIS) we are interested in finding an independent set S such that $w(S) \geq w(S')$ for all independent sets S' . This problem is known to be NP-hard in general graphs.

Our approach is based on the concept of dynamic programming over tree decompositions. A tree decomposition is a cover of the vertex set satisfying certain structural properties, amongst which is the requirement that the elements of the cover form a tree. The tree structure facilitates a bottom-up computation where at any given point the problem is restricted to the induced subgraph of the current cover element. This leads to a parameterized algorithm, where the running time is exponential in the width of the decomposition but polynomial in the size of the input graph.

2 Preliminaries

The material in this section is adapted from Cygan et al. [4].

2.1 Separators

Definition 2.1. A *separation* of G is a pair (A, B) , where $A, B \subseteq V(G)$, such that $A \cup B = V(G)$, and for which there are no edges from $A \setminus B$ to $B \setminus A$. The intersection $A \cap B$ is then called the *separator*.

For two independent sets S of $G[A]$ and S' of $G[B]$ we say that they *agree* on the separator if $S \cap (A \cap B) = S' \cap (A \cap B)$. When they agree, their union $S \cup S'$ is an independent set of G .

Figure 1 illustrates a separation of a graph, its separator $A \cap B$, and two independent sets that agree on the separator.

2.2 Tree Decompositions

Definition 2.2. A *tree decomposition* TD of a graph G is a tuple $(T, \{X_t\}_{t \in N(T)})$, where T is a tree, the vertices of which are referred to as *nodes* and the set of which is denoted by $N(T)$, together with a family of *bags* $\{X_t\}_{t \in N(T)}$, such that for every node $t \in N(T)$, we have $X_t \subseteq V(G)$, satisfying the following requirements:

Author's address: Bjarn de Jong, bjarrn@gmail.com.

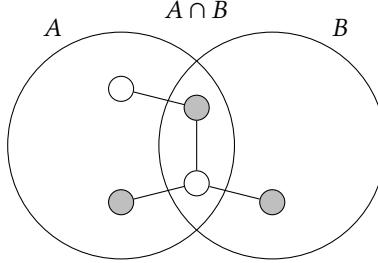


Fig. 1. Separation and separator

- *Vertex coverage*: For every $u \in V(G)$, there exists a node $t \in N(T)$ such that $u \in X_t$.
- *Edge coverage*: For every edge $\{u, v\} \in E(G)$, there exists a node $t \in N(T)$ such that $\{u, v\} \subseteq X_t$.
- *Running intersection property*: For every $u \in V(G)$, the induced subgraph $T[\{t \in N(T) \mid u \in X_t\}]$ is a subtree of T .

The *width* of a tree decomposition TD is defined as $\text{wd}(TD) = \max\{|X_t| \mid t \in N(TD)\} - 1$, and the *treewidth* of a graph G is defined as $\text{tw}(G) = \min\{\text{wd}(TD) \mid TD \text{ is a tree decomposition of } G\}$.

LEMMA 2.3. *For every pair of adjacent nodes $t, t' \in N(T)$, we have that $X_t \cap X_{t'}$ is a separator of the graph G .*

2.3 Rooted Tree Decompositions

We choose an arbitrary node as the *root* for the tree and will henceforth only work with *rooted tree decompositions*. The dependency relation induces a partial order on the nodes of the tree, where for $t, t' \in N(T)$ we write $t' \leq t$ if and only if t' is a descendent of t . We will consider separations of the form

$$(A, B) = (G[\bigcup_{a \leq t'} X_a], G[\bigcup_{b \notin t'} X_b])$$

with separator $X_{t'} \cap X_t$, where t is the parent of t' . For a node t , we write G_t for the induced subgraph $G[\bigcup_{t' \leq t} X_{t'}]$ and V_t for $V(G_t)$.

Definition 2.4. Given a rooted tree decomposition TD of a graph G , we define for every node t and every independent set $S \subseteq X_t$

$$c[t, S] = \max \left\{ w(\widehat{S}) \mid \begin{array}{l} S \subseteq \widehat{S} \subseteq V_t, \\ \widehat{S} \text{ is an independent set,} \\ X_t \cap \widehat{S} = S \end{array} \right\}.$$

The aim will be to traverse up the tree and for each node $t \in N(T)$ and each independent set $S \subseteq X_t$ to compute $c[t, S]$. Finally, upon reaching the root, the weight of a MWIS will be

$$\max\{c[\text{root}, S] \mid S \subseteq X_{\text{root}}, S \text{ independent set}\}.$$

2.4 Nice Tree Decompositions

Definition 2.5. A rooted tree decomposition is called *nice* if it is a binary tree satisfying the following conditions:

- The *leaf nodes* and the *root* satisfy:
 - For every leaf node l , we have $X_l = \emptyset$.
 - For the root node r , we also have $X_r = \emptyset$.

- Every *non-leaf node* is one of the following three types:
 - *Introduce node*: a node t with exactly one child t' , such that $X_t = X_{t'} \sqcup \{v\}$ for some vertex $v \in V$.
 - *Forget node*: a node t with exactly one child t' , such that $X_t \sqcup \{v\} = X_{t'}$ for some vertex $v \in V$.
 - *Join node*: a node t with exactly two children t_1 and t_2 , such that $X_t = X_{t_1} = X_{t_2}$.

Computation of $c[\cdot, \cdot]$ in a Nice Tree Decomposition For a nice tree decomposition, the values $c[\cdot, \cdot]$ can be computed in a bottom-up manner as follows:

- For every leaf node l :
 - We have $X_l = \emptyset$, so:
- $$c[l, \emptyset] = 0.$$
- For non-leaf nodes:
 - Introduce node: Let t be an introduce node where vertex v is introduced, and let t' be its child. Then:

$$c[t, S] = \begin{cases} c[t', S], & \text{if } v \notin S, \\ c[t', S \setminus \{v\}] + w(v), & \text{if } v \in S. \end{cases}$$

- Forget node: Let t be a forget node where vertex v is forgotten, and let t' be its child. Then:

$$c[t, S] = \begin{cases} \max\{c[t', S], c[t', S \sqcup \{v\}]\}, & \text{if } S \sqcup \{v\} \text{ is independent,} \\ c[t', S], & \text{otherwise.} \end{cases}$$

- Join node: Let t be a join node with children t_1 and t_2 . Then:

$$c[t, S] = c[t_1, S] + c[t_2, S] - w(S).$$

3 Algorithm

Instead of working with nice tree decompositions we will work with arbitrary rooted tree decompositions and borrow concepts from nice tree decompositions. The overall traversal of the tree decomposition is implemented using a depth-first traversal. During the traversal, we compute dynamic programming tables. Several vectors are used to represent the necessary information at each node. Most notably, we represent all independent sets local to a bag using natural numbers, where each number encodes a subset via its binary representation. These encoded sets are manipulated efficiently using bitwise operations and basic arithmetic.

3.1 Input

As input the program takes two arguments, a .td file representing the tree decomposition and a .graph file representing the vertex-weighted graph.

.graph file The .graph file consist of a first line $n \ m \ 10$, where n is the number of vertices, m is the number of edges and 10 indicates that vertices have weights. The vertices are 1-indexed and thus make up the range $[1, n]$. This is then followed by n lines of the form $w(i) \ N(i)$, where $w(i)$ is the weight of the vertex i and where $N(i)$ is the neighborhood of vertex i listed in ascending order.

.td file The .td file format is as specified in the PACE 2017 Treewidth track [3]. Additionally, we require that the bags are sorted, that is, X_i comes before X_j if $i < j$ and the contents of the bags are sorted in ascending order. We also require the edge list to be sorted in ascending order where $(i < j) < (k < l)$ if $i < k$ or $i = k$ and $j < l$.

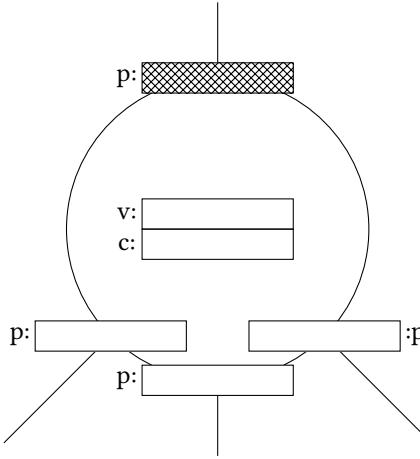


Fig. 2. Storage of the data structure locally at a node t .

3.2 Data Structure

Encoding Per node t , enumerate all independent sets local to the bag X_t encoded as integers in the range $[0, 2^{|X_t|} - 1]$. The encoding is given such that the i -th bit in the binary representation of an independent set is set to 1 if and only if the i -th vertex of the bag is included in the independent set.

Storage Let t be a node with bag X_t . Then the data structure used locally for t consists of several vectors, all of which have the same length. The vector v contains all the independent sets $S \subseteq X_t$ encoded as integers in ascending order. The vector c is to contain the values $c[t, \cdot]$ from Definition 2.4. For each child t' of t , we have a vector p to track how each independent set at t relates to a corresponding independent set at t' . Figure 2 illustrates the local data structure at a node t .

3.3 Computation

As mentioned above, instead of working with a nice tree decomposition, we work with an arbitrary rooted tree decomposition. The depth-first traversal determines the order in which computations occur. Finishing a node involves three steps: initializing the node if it is a leaf, updating the current node's information to account for multiple children, and pushing the relevant information from the current node to its parent.

Updating to the Parent Node To update information from a current node t' to its parent t , we traverse a *virtual path* consisting of forget nodes until we reach the intersection of the bags X_t and $X_{t'}$, followed by a virtual path of introduce nodes until X_t is reached. In particular, assuming $X_t \not\subseteq X_{t'}$ and $X_{t'} \not\subseteq X_t$, the virtual node corresponding to the intersection $X_t \cap X_{t'}$ acts as a forget node, while t itself is considered an introduce node in this virtual path (see Figure 3).

Computing the vectors c upward follows the approach used in nice tree decompositions, except that $c[t, S]$ is initialized to zero, and the contributions from the virtual path at t are added rather than assigned directly. Traversing this virtual path consists of *steps*, where a step either forgets a vertex at the i -th position from the bag or introduces a vertex at the i -th position to the bag. Figure 4 illustrates the behavior of these virtual steps for forgetting and introducing a vertex. For the purpose of computing our vectors v , c and p from the data structure, whilst forgetting the i -th vertex v from the bag, every independent S not containing v is maintained and we only need to consider whether or not $S \sqcup \{v\}$ was also an independent set. In the case of introducing an i -th

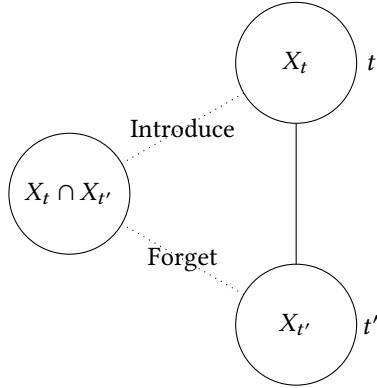
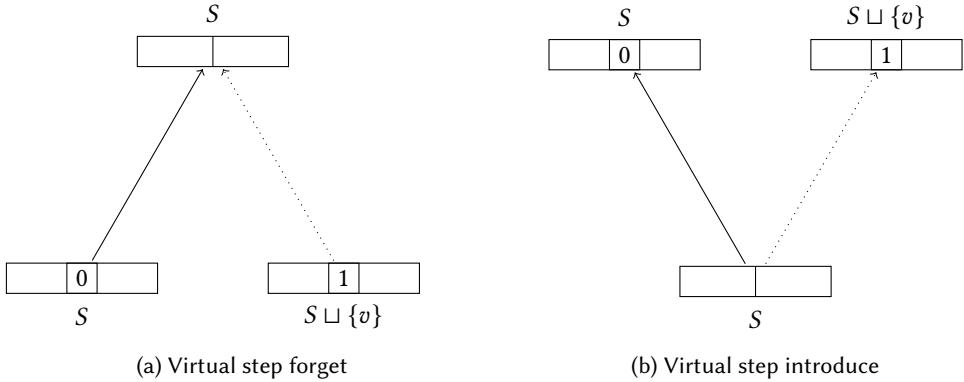
Fig. 3. Virtual path from node t' to its parent t 

Fig. 4. Treatment of independent sets during the traversal of the virtual path using the binary encodings.

vertex v to the bag, every independent set S is maintained, and we only need to consider whether or not such an independent set gives rise to an additional independent set $S \sqcup \{v\}$.

Updating the Current Node Since our tree decomposition nodes may have arbitrary degrees (no join nodes), we need to subtract the contribution of a set S from $c[t, S]$ the appropriate number of times. Unless t is the root (in which case the factor is 1 instead of 2), this update can be expressed as

$$c[t, S] = c[t, S] - (\deg(t) - 2) \cdot w(S),$$

where $\deg(t)$ denotes the degree of node t and $w(S)$ the weight of set S .

Initializing Leaf Nodes Leaf initialization is a special case of updating to the parent, where we update from a virtual empty node to the leaf node.

3.4 Memory

For a node t with bag X_t , the vectors used by the algorithm can be quite large, containing anywhere between 1 and $2^{|X_t|}$ entries. Since storing many such large vectors simultaneously can put significant pressure on memory resources, we consider their lifespans and how they are managed.

The p vectors The lifespan of the p vectors depends on whether the actual maximum weight independent set is computed, or if only its weight is required. The p vectors are optional and only computed alongside the c vectors if the maximum weight independent set itself is needed. Once computed, they must persist for the entire duration of the program, as they are used during a root-to-leaf traversal to reconstruct the solution. Because the p vectors are only needed when reading out the final solution, we manage them by writing to and reading from disk using the Zstandard compression library to reduce memory usage.

The v vectors The v vectors are only required for the current node during the depth-first traversal.

The c vectors The c vectors are required for the current node, but must also persist for all *active* nodes. A node t is considered *active* if it has at least one finished child and at least one unfinished child. Because the number of active nodes can grow substantially for large tree decompositions, we provide an option to store these c vectors on disk as well, again using the Zstandard library for compression.

3.5 Output

The algorithm outputs the weight of a maximum weight independent set on the first line, followed by the time taken to compute this value in seconds on the second line. If the actual independent set is requested via an optional flag, an additional line is printed, listing the vertices included in the set in ascending order.

3.6 Running Time Analysis

Let G be a graph and let TD be a rooted tree decomposition of G with $k = \text{wd}(TD) + 1$. Then the algorithm runs within $O(n \cdot k \cdot 2k \cdot 2^k)$, where 2^k is an upperbound on the size of a DP table, $2k$ is an upperbound on the length of a virtual path, k is an upperbound on the effort required to check if a set is independent and n is an upperbound on the number of nodes in the tree decomposition.

4 Results

The primary bottleneck of our algorithm is the exponential number of independent sets in the bags of the tree decomposition. Although the theoretical upper bound for the number of independent sets encountered in a bag of a tree decomposition TD with $k = \text{width}(TD) + 1$ is 2^k , the actual number of sets encountered in practice is often much smaller. We can reliably solve the MWIS problem for all graphs with a tree decomposition of maximal bag size below 32, and provide support for decompositions with maximal bag size below 128, where the latter corresponds to the maximum size our encoding can handle.

4.1 Tree Decomposition Method

All tree decompositions used in this work were computed using the heuristic algorithm provided in the PACE 2017 Track A repository by Tamaki, Ohtsuka, Sato, and Makii [9], which incorporates the exact method described by Tamaki [8] as one of its components.

4.2 System Configuration for Performance Evaluation

All experiments and performance measurements were conducted on the following hardware and software environment:

- CPU: Intel Core i5-2400S @ 2.50 GHz, 4 cores (no hyper-threading), x86_64 architecture
- RAM: 23 GiB
- Operating System: Ubuntu 24.04.2 LTS, Linux kernel 6.11.0-26-generic

- Compiler: `g++` with flags:
 - `-O3`: Enables aggressive compiler optimizations
 - `-march=native`: Targets the host CPU with all available instruction sets.
 - `-std=c++20`: Uses the C++20 standard

4.3 Datasets

The benchmark set used in this work includes graphs drawn from standard datasets commonly used in MWIS research. These instances originate from publicly available sources, including OpenStreetMap (OSM), the Stanford Large Network Dataset Repository (SNAP), the SuiteSparse Matrix Collection (SSMC), and various 3D mesh and finite element datasets. We reused these datasets similar as to other works [2, 5, 7], all of which evaluate solvers for graph problems on subsets of the same graph families. While not distributed as a single unified benchmark suite, these graphs can be either reconstructed using the data sources and scripts referenced in the cited works or are shared amongst the community. To ensure consistency and comparability, we follow the instance naming conventions established in prior literature where applicable.

4.4 Results on the Dataset

Our dataset consists of 213 graphs grouped into five categories:

- (mesh) 14 graphs from well-known triangle meshes.
- (fe) 5 3D meshes derived from simulations using the finite element method.
- (snap) 40 graphs from the Stanford Large Network Dataset Repository.
- (ssmc) 6 graphs from the SuiteSparse Matrix Collection.
- (osm) 148 conflict graphs generated from OpenStreetMap data, following Barth et al. [1].

We excluded the `snap` graphs from the following analysis due to difficulties in computing suitable tree decompositions. For the remaining four groups, detailed tables presenting the tree decomposition widths and the corresponding results of our MWIS solver on the benchmark set are provided in the appendix, see Appendix A and Appendix B. Only the `osm` group contained graphs with widths within our guaranteed solvable range. In fact, 112 of the 148 `osm` graphs had widths supported by our datatype, and all of these were solvable by our algorithm.

Figure 5 plots for the `osm` graphs the maximum bag size occurring in the computed tree decompositions against the logarithm of the maximal number of independent sets in a bag. We use a grayscale color map to indicate the size of the corresponding bag, since the bag containing the maximal number of sets does not always correspond to the one with size closest to the decomposition's width.

The information in Appendix A and Appendix B also enables comparison against LearnAndReduce, a separate MWIS solver by Großmann et al [6]; see also the accompanying paper [5] for further details. While our solver could not solve any instances that LearnAndReduce could not fully reduce, we were able to successfully solve several reduced instances produced by LearnAndReduce by subsequently applying our algorithm. These additional instances and results are presented in Appendix C.

5 Future Work

Currently, the solver supports tree decompositions with maximum bag sizes less than 128. Since all `osm` graphs within this bound were solvable by the algorithm, extending support to maximum bag sizes below 255 might result in additional solved instances.

The tree decompositions used were computed using heuristics from [9]. These heuristics produce generally good decompositions, but were not designed with the specific needs of this MWIS solver in mind. Although our focus was on implementing the MWIS solver for given tree decompositions,

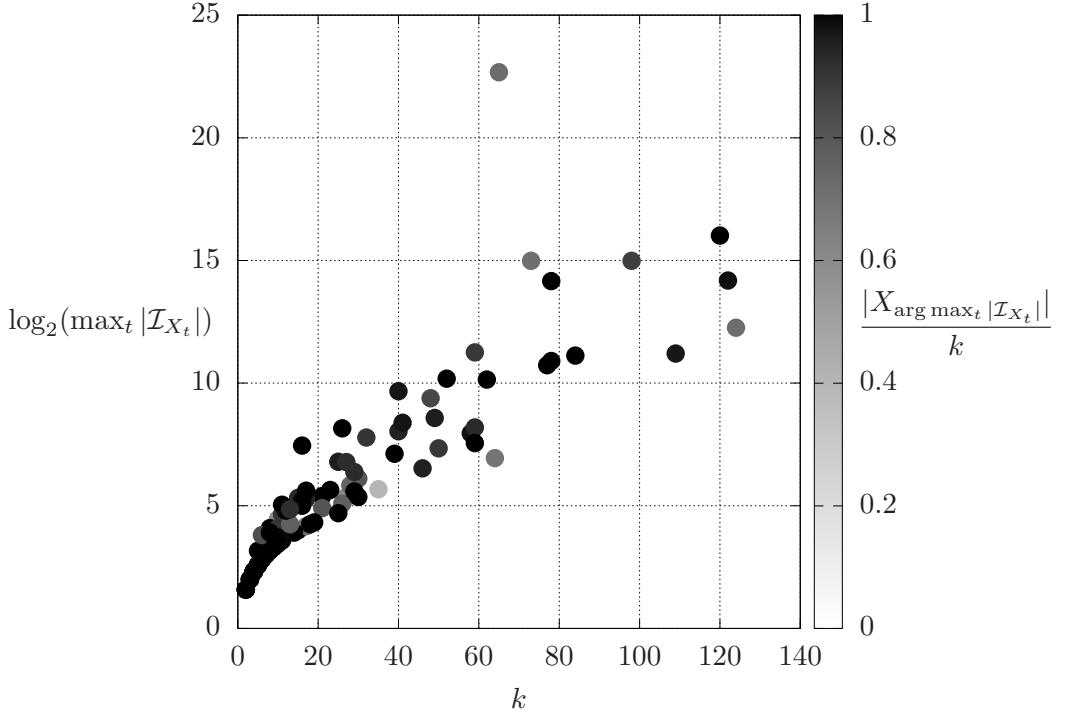


Fig. 5. All plotted points correspond to OSM instances for which we obtained a tree decomposition with maximal bag size k less than 128, all of which are solvable by the algorithm described in this paper. The scatter plot shows the relationship between k and the base-2 logarithm of the maximum value of $|\mathcal{I}_{X_t}|$ taken over all bags in the tree decomposition, where \mathcal{I}_{X_t} denotes the set of independent sets of the induced subgraph of the bag X_t . The grayscale shading indicates the ratio between the size of the bag where this maximum occurs and k .

it would likely be beneficial to develop a tree decomposition algorithm that specifically minimizes the maximal number of independent sets occurring in a bag. If such an approach is feasible, it could significantly improve the efficiency and scalability of the solver.

In addition, extending the exact solver towards a heuristic approach could improve scalability.

References

- [1] Lukas Barth, Benjamin Niedermann, Martin Nöllenburg, and Darren Strash. Temporal map labeling: a new unified framework with experiments. In *Proceedings of the 24th ACM SIGSPATIAL International Conference on Advances in Geographic Information Systems*, SIGSPATIAL '16, New York, NY, USA, 2016. Association for Computing Machinery. ISBN 9781450345897. doi: 10.1145/2996913.2996957. URL <https://doi.org/10.1145/2996913.2996957>.
- [2] Shaowei Cai, Wenyi Hou, Jinkun Lin, and Yuanjie Li. Improving Local Search for Minimum Weight Vertex Cover by Dynamic Strategies. In *Proceedings of the Twenty-Seventh International Joint Conference on Artificial Intelligence, IJCAI-18*, pages 1412–1418. International Joint Conferences on Artificial Intelligence Organization, 7 2018. doi: 10.24963/ijcai.2018/196. URL <https://doi.org/10.24963/ijcai.2018/196>.
- [3] PACE Challenge Committee. PACE Challenge 2017: Treewidth Track. <https://pacechallenge.org/2017/treewidth/>, 2017. Accessed: 2025-04-28.
- [4] Marek Cygan, Fedor V. Fomin, Łukasz Kowalik, Daniel Lokshtanov, Dániel Marx, Marcin Pilipczuk, Michał Pilipczuk, and Saket Saurabh. *Parameterized Algorithms*, chapter 7, pages 157–166. Springer, Cham, May 2015. ISBN 978-3-319-21275-3. doi: 10.1007/978-3-319-21275-3. URL <https://doi.org/10.1007/978-3-319-21275-3>.
- [5] Ernestine Großmann, Kenneth Langedal, and Christian Schulz. Accelerating Reductions Using Graph Neural Networks and a New Concurrent Local Search for the Maximum Weight Independent Set Problem. Mar 2025. doi: 10.48550/arXiv.2412.14198. URL <https://arxiv.org/abs/2412.14198v2>.
- [6] Ernestine Großmann, Kenneth Langedal, and Christian Schulz. KarlsruheMIS/LearnAndReduce: v1.0 - ACDA2025, Apr 2025. URL <https://doi.org/10.5281/zenodo.15229975>.
- [7] Sebastian Lamm, Christian Schulz, Darren Strash, Robert Williger, and Huashuo Zhang. Exactly Solving the Maximum Weight Independent Set Problem on Large Real-World Graphs. In *2019 Proceedings of the Twenty-First Workshop on Algorithm Engineering and Experiments (ALENEX)*, pages 144–158. SIAM, 2019. doi: 10.1137/1.9781611975499.12. URL <https://pubs.siam.org/doi/abs/10.1137/1.9781611975499.12>.
- [8] Hisao Tamaki. Positive-instance driven dynamic programming for treewidth. *arXiv preprint arXiv:1704.05286*, 2017. doi: 10.48550/arXiv.1704.05286. URL <https://arxiv.org/abs/1704.05286>.
- [9] Hisao Tamaki, Hiromu Ohtsuka, Takuto Sato, and Keitaro Makii. PACE2017-TrackA. <https://github.com/TCS-Meiji/PACE2017-TrackA>, 2017. Accessed: 2025-04-28.

A Results on OSM

Table 1. Results on OSM instances (part 1 of 4). Each row corresponds to a graph instance derived from OpenStreetMap data. The columns report the number of vertices n , edges m , decomposition width plus one k (i.e., the size of the largest bag), and number of bags N in the computed tree decomposition. The columns w and t respectively show the weight of the computed maximum weight independent set and the runtime (in seconds) of our solver introduced in this work. A dash (-) indicates that no result could be obtained by our solver. Gray shading is used to indicate that an instance is known to be solvable by LearnAndReduce, independent of our solver.

OSM instance	n	m	k	N	w	t
alabama-AM1	320	581	10	206	167 588	0.0006
alabama-AM2	1 164	19 386	73	447	174 309	0.1071
alabama-AM3	3 504	309 664	436	742	-	-
alaska-AM1	31	31	4	13	20 266	<0.0001
alaska-AM2	54	156	11	20	20 900	<0.0001
alaska-AM3	86	475	21	36	22 325	0.0002
arkansas-AM1	26	19	4	11	17 702	<0.0001
arkansas-AM2	55	233	15	20	20 771	0.0001
arkansas-AM3	103	1 376	40	28	20 935	0.0007
california-AM1	77	130	8	38	46 537	<0.0001
california-AM2	231	3 074	40	87	47 153	0.0029
california-AM3	587	27 536	189	144	-	-
canada-AM1	189	240	8	92	78 466	0.0002
canada-AM2	449	2 947	30	153	81 799	0.0008
canada-AM3	943	20 241	98	281	86 018	0.1397
colorado-AM1	128	232	7	56	50 507	0.0001
colorado-AM2	283	2 026	27	85	52 172	0.0004
colorado-AM3	538	8 365	62	164	54 741	0.0040
connecticut-AM1	87	96	5	44	55 131	<0.0001
connecticut-AM2	211	975	17	65	56 058	0.0002
connecticut-AM3	367	3 769	46	118	57 650	0.0007
delaware-AM1	2	1	2	1	1 060	<0.0001
delaware-AM2	3	3	3	1	1 060	<0.0001
delaware-AM3	5	9	4	2	1 060	<0.0001
district-of-columbia-AM1	2 500	24 651	65	1 378	196 475	9.8375
district-of-columbia-AM2	13 597	1 609 795	857	2 644	-	-
district-of-columbia-AM3	46 221	27 729 137	5 106	5 126	-	-
florida-AM1	475	1 277	18	230	225 655	0.0005
florida-AM2	1 254	16 936	58	387	230 595	0.0037
florida-AM3	2 985	154 043	230	783	-	-
georgia-AM1	294	434	8	144	205 068	0.0003
georgia-AM2	746	7 753	64	256	216 346	0.0019
georgia-AM3	1 680	74 126	267	480	-	-
greenland-AM1	77	341	16	46	9 328	0.0002
greenland-AM2	686	50 218	210	152	-	-
greenland-AM3	4 986	3 652 361	2 830	357	-	-
hawaii-AM1	411	1 423	15	227	113 792	0.0009

Table 2. Results on OSM instances (part 2 of 4). Continuation of Table 1.

OSM instance	<i>n</i>	<i>m</i>	<i>k</i>	<i>N</i>	<i>w</i>	<i>t</i>
hawaii-AM2	2 875	265 158	320	554	-	-
hawaii-AM3	28 006	49 444 921	9 208	1 426	-	-
idaho-AM1	136	208	11	63	70 623	0.0002
idaho-AM2	552	35 221	186	123	-	-
idaho-AM3	4 064	3 924 080	2 957	272	-	-
illinois-AM1	113	202	10	58	54 678	0.0002
illinois-AM2	261	2 138	32	101	55 496	0.0011
indiana-AM1	2	1	2	1	1 146	<0.0001
indiana-AM2	2	1	2	1	1 146	<0.0001
indiana-AM3	4	6	4	1	1 146	<0.0001
iowa-AM1	90	164	7	38	47 907	<0.0001
iowa-AM2	155	954	21	49	47 984	0.0003
kansas-AM1	190	400	11	92	84 449	0.0002
kansas-AM2	602	16 474	120	183	85 942	0.1932
kansas-AM3	2 732	806 912	1 166	347	-	-
kentucky-AM1	381	2 402	39	184	91 897	0.0010
kentucky-AM2	2 453	643 428	914	398	-	-
kentucky-AM3	19 095	59 533 630	12 625	966	-	-
louisiana-AM1	157	181	6	93	51 446	0.0002
louisiana-AM2	436	3 111	28	154	55 127	0.0008
louisiana-AM3	1 162	37 077	137	299	-	-
maine-AM1	38	29	4	19	24 921	<0.0001
maine-AM2	81	243	10	25	26 208	<0.0001
maine-AM3	143	850	17	49	26 734	0.0003
maryland-AM1	104	216	10	47	43 930	0.0001
maryland-AM2	316	4 715	50	99	45 300	0.0010
maryland-AM3	1 018	95 415	281	168	-	-
massachusetts-AM1	413	1 089	16	237	136 695	0.0009
massachusetts-AM2	1 339	35 449	160	440	-	-
massachusetts-AM3	3 703	551 491	1 014	709	-	-
mexico-AM1	175	358	10	78	90 599	0.0002
mexico-AM2	516	9 411	59	172	94 834	0.0025
mexico-AM3	1 096	47 131	170	286	-	-
michigan-AM1	133	112	4	69	51 076	0.0001
michigan-AM2	241	750	14	87	51 928	0.0002
michigan-AM3	376	2 459	29	138	52 674	0.0009
minnesota-AM1	86	136	5	52	28 692	0.0002

Table 3. Results on OSM instances (part 3 of 4). Continuation of Table 1.

OSM instance	<i>n</i>	<i>m</i>	<i>k</i>	<i>N</i>	<i>w</i>	<i>t</i>
minnesota-AM2	253	2 580	35	100	30 251	0.0007
minnesota-AM3	683	34 188	182	171	-	-
mississippi-AM1	74	60	5	39	32 273	<0.0001
mississippi-AM2	151	366	10	48	33 187	0.0001
mississippi-AM3	242	1 116	16	90	33 318	0.0003
missouri-AM1	10	6	2	6	7 928	<0.0001
missouri-AM2	13	12	3	6	7 928	<0.0001
missouri-AM3	17	24	4	8	7 928	<0.0001
montana-AM1	109	194	9	49	55 348	0.0001
montana-AM2	307	5 154	59	100	56 068	0.0026
montana-AM3	837	69 293	297	198	-	-
nebraska-AM1	40	46	7	17	24 345	<0.0001
nebraska-AM2	93	734	26	31	26 680	0.0002
nebraska-AM3	145	2 168	49	47	27 214	0.0014
nevada-AM1	89	93	6	46	45 761	0.0001
nevada-AM2	242	1 531	25	81	47 068	0.0004
nevada-AM3	569	15 016	109	156	52 036	0.0156
new-hampshire-AM1	195	302	7	98	108 186	0.0002
new-hampshire-AM2	514	3 369	25	181	110 621	0.0009
new-hampshire-AM3	1 107	18 021	59	367	116 060	0.0151
new-jersey-AM1	4	6	4	1	256	<0.0001
new-jersey-AM2	4	6	4	1	256	<0.0001
new-jersey-AM3	4	6	4	1	256	<0.0001
new-mexico-AM1	3	3	3	1	182	<0.0001
new-mexico-AM2	3	3	3	1	182	<0.0001
new-mexico-AM3	3	3	3	1	182	<0.0001
new-york-AM1	42	118	11	20	13 187	<0.0001
new-york-AM2	224	6 399	78	72	14 330	0.0445
new-york-AM3	837	88 728	329	137	-	-
north-carolina-AM1	93	150	7	46	45 254	0.0001
north-carolina-AM2	398	10 116	78	89	46 896	0.0103
north-carolina-AM3	1 557	236 739	353	223	-	-
ohio-AM1	78	96	6	42	50 964	<0.0001
ohio-AM2	211	1 815	30	67	51 289	0.0004
ohio-AM3	482	11 376	84	151	52 634	0.0107
oregon-AM1	381	996	17	203	161 298	0.0005
oregon-AM2	1 325	57 517	202	361	-	-

Table 4. Results on OSM instances (part 4 of 4). Continuation of Table 1.

OSM instance	<i>n</i>	<i>m</i>	<i>k</i>	<i>N</i>	<i>w</i>	<i>t</i>
oregon-AM3	5 588	2 912 701	2 332	705	-	-
pennsylvania-AM1	193	276	8	99	133 914	0.0002
pennsylvania-AM2	521	3 812	27	173	138 413	0.0010
pennsylvania-AM3	1 148	26 464	124	359	143 870	0.0452
puerto-rico-AM1	60	63	5	30	29 802	<0.0001
puerto-rico-AM2	165	1 285	29	40	32 921	0.0003
puerto-rico-AM3	494	26 926	175	84	-	-
rhode-island-AM1	455	1 973	26	269	171 224	0.0012
rhode-island-AM2	2 866	295 488	402	692	-	-
rhode-island-AM3	15 124	12 622 219	5 772	1 380	-	-
south-carolina-AM1	75	69	5	42	50 033	<0.0001
south-carolina-AM2	165	713	16	64	51 446	0.0003
south-carolina-AM3	317	4 508	48	103	52 087	0.0025
tennessee-AM1	49	39	4	23	29 569	<0.0001
tennessee-AM2	100	418	19	30	31 567	0.0001
tennessee-AM3	212	3 215	52	55	32 276	0.0062
utah-AM1	230	309	6	125	87 856	0.0002
utah-AM2	589	4 692	41	212	95 087	0.0018
utah-AM3	1 339	42 872	181	386	-	-
vermont-AM1	128	418	13	65	55 884	0.0002
vermont-AM2	766	37 607	135	171	-	-
vermont-AM3	3 436	1 136 164	1 012	319	-	-
virginia-AM1	570	1 480	12	287	280 936	0.0007
virginia-AM2	2 279	60 040	122	662	295 867	0.1249
virginia-AM3	6 185	665 903	579	1 244	-	-
washington-AM1	713	2 316	23	355	296 653	0.0019
washington-AM2	3 025	152 449	255	759	-	-
washington-AM3	10 022	2 346 213	1 470	1 606	-	-
west-virginia-AM1	65	150	8	33	42 868	<0.0001
west-virginia-AM2	317	8 328	77	78	45 923	0.0053
west-virginia-AM3	1 185	125 620	335	206	-	-
wisconsin-AM1	54	51	4	24	44 608	<0.0001
wisconsin-AM2	89	219	10	27	44 651	<0.0001
wisconsin-AM3	136	588	13	56	47 904	0.0004
wyoming-AM1	7	11	5	2	4 568	<0.0001
wyoming-AM2	8	16	6	2	4 568	<0.0001
wyoming-AM3	12	42	9	3	4 568	<0.0001

B Results on fe, mesh and ssmc

Table 5. Results on fe, mesh, and ssmc instance groups. Each row corresponds to a graph instance derived from the specified group. The columns report the number of vertices n , edges m , decomposition width plus one k (i.e., the size of the largest bag), and number of bags N in the computed tree decomposition. Rows shaded in gray indicate instances known to be solvable by LearnAndReduce. Since our solver did not produce solutions on any of these instances, the weight and runtime columns are omitted.

	Instance	n	m	k	N
fe	body-uniform	45 087	163 734	236	30 563
	ocean-uniform	143 437	409 593	143 398	40
	pwt-uniform	36 519	144 794	214	18 943
	rotor-uniform	99 617	662 431	99 617	1
	sphere-uniform	16 386	49 152	414	9 805
mesh	blob-uniform	16 068	24 102	166	13 262
	buddha-uniform	1 087 716	1 631 574	1 087 716	1
	bunny-uniform	68 790	103 017	349	56 333
	cow-uniform	5 036	7 366	43	4 301
	dragonsub-uniform	600 000	900 000	600 000	1
	dragon-uniform	150 000	225 000	232	125 637
	ecat-uniform	684 496	1 026 744	2 006	555 191
	face-uniform	22 871	34 054	129	19 126
	fandisk-uniform	8 634	12 818	97	7 159
	feline-uniform	41 262	61 893	141	34 000
	gameguy-uniform	42 623	63 850	241	35 759
	gargoyle-uniform	20 000	30 000	149	16 682
	turtle-uniform	267 534	401 178	1 278	220 527
	venus-uniform	5 672	8 508	75	4 750
ssmc	ca2010	710 145	1 744 683	395	608 354
	fl2010	484 481	1 173 147	297	422 722
	ga2010	291 086	709 028	493	250 917
	il2010	451 554	1 082 232	551	376 861
	nh2010	48 837	117 275	142	42 850
	ri2010	25 181	62 875	78	21 334

C Results on Partially Reduced Instances

Table 6. Results on partially reduced instances provided by LearnAndReduce. Each row corresponds to a graph instance that was preprocessed by LearnAndReduce, after which we applied our solver to the reduced instance. The columns report the number of vertices n , edges m , decomposition width plus one k (i.e., the size of the largest bag), and number of bags N in the computed tree decomposition of the reduced instance. The columns w and t respectively show the weight of the computed maximum weight independent set and the runtime (in seconds) of our solver on these reduced instances. A dash (-) indicates that no result could be obtained.

	Reduced instance	n	m	k	N	w	t
fe	body-uniform	396	3 571	43	191	7 713	16.7516
	pwt-uniform	15 274	107 248	116	7 294	-	-
	rotor-uniform	89 294	690 242	2 128	40 946	-	-
osm	alabama-AM3	815	61 637	239	157	-	-
	california-AM3	359	19 705	159	65	-	-
	district-of-columbia-AM2	2 361	274 528	584	465	-	-
	district-of-columbia-AM3	26 728	13 406 329	3 481	3 010	-	-
	florida-AM3	658	41 043	162	107	-	-
	georgia-AM3	481	31 388	191	103	-	-
	greenland-AM3	3 402	2 141 096	2 296	241	-	-
	hawaii-AM3	22 849	38 081 016	8 031	911	-	-
	idaho-AM3	2 920	2 413 139	2 386	130	-	-
	kansas-AM3	1 209	292 997	753	117	-	-
	kentucky-AM3	16 714	53 762 869	12 113	712	-	-
	massachusetts-AM3	1 626	340 652	837	216	-	-
	mexico-AM3	447	20 443	105	118	11 196	1.6672
	montana-AM3	382	39 859	231	63	-	-
	new-york-AM3	526	50 572	266	98	-	-
	north-carolina-AM3	997	133 106	293	135	-	-
	oregon-AM3	3 159	1 817 045	1 950	214	-	-
	rhode-island-AM2	498	46 183	187	147	-	-
	rhode-island-AM3	12 189	11 225 470	5 349	866	-	-
	vermont-AM3	1 946	526 463	790	167	-	-
	virginia-AM3	2 385	293 068	408	359	-	-
	washington-AM3	6 671	1 955 568	1 388	687	-	-
	west-virginia-AM3	864	98 399	325	141	-	-
snap	as-skitter-uniform	1 996	27 527	222	993	-	-
	loc-gowalla_edges	304	2 938	104	184	-	-
	soc-LiveJournal1-uniform	1 360	21 602	137	780	-	-
	soc-p.-rel.-uniform	782 844	14 236 948	782 844	1	-	-
	web-NotreDame	75	1 178	37	24	2 728	0.0337
	web-NotreDame-uniform	97	1 357	38	43	3 323	0.9553